

Topic Modeling

Jonathan Neimann

2024-11-01

Note: I have set the seed on the lda but i am still getting different results in my clusters. the code should work fine but the results i write about might be different than what you are seeing on the pdf

```
movies = read.csv("movie_plots_with_genres.csv")
plots_by_word = movies %>% unnest_tokens(word, Plot)
plot_word_counts = plots_by_word %>% anti_join(stop_words) %>%
count(Movie.Name, word, sort=TRUE)

## Joining with `by = join_by(word)`

data("freq_first_names")
first_names = tolower(freq_first_names$Name)
plot_word_counts = plot_word_counts %>% filter(!(word %in% first_names))

plots_dtm = plot_word_counts %>% cast_dtm(Movie.Name, word, n)
```


Create a screen plot to optimize the number of topics we should use

```
# Define a range of topic numbers (e.g., 2 to 30)
k_values <- seq(2, 30, by = 2)

# Initialize a data frame to store the results
perplexity_results <- data.frame(k = integer(), perplexity = numeric())

# Loop over each k and fit an LDA model, storing the perplexity
for (k in k_values) {
  lda_model <- LDA(plots_dtm, k = k, control = list(seed = 123))
  perplexity_val <- perplexity(lda_model, plots_dtm)
  perplexity_results <- rbind(perplexity_results, data.frame(k = k,
perplexity = perplexity_val))
}

# Plot the perplexity against the number of topics (k)
library(ggplot2)
ggplot(perplexity_results, aes(x = k, y = perplexity)) +
  geom_line() +
  geom_point() +
  labs(title = "Scree Plot for Optimal Number of Topics",
       x = "Number of Topics (k)",
       y = "Perplexity") +
  theme_minimal()
```



scree plot does not seem to have an “elbow” that determines our optimal number of topics, so this can be up to the user’s discretion. I am going to go with 20.

LDA with 20 topics

```
plots_lda = LDA(plots_dtm, k = 20, control = list(seed=123))
```

Now we need to retrieve the gammas from this lda, which represent the topics.

#retrieving gammas

```
betas = tidy(plots_lda, matrix = "beta")
betas_wider = betas %>% pivot_wider(names_from = topic, values_from = beta)

plots_gamma = tidy(plots_lda, matrix = "gamma")
plots_gamma_wider = plots_gamma %>% pivot_wider(names_from = topic,
                                              values_from = gamma
                                              )
```


We can now take these clusters and assign them to each movie in the dataframe

```
#drop duplicate movie names to match the rows
movies <- movies %>%
  distinct(Movie.Name, .keep_all = TRUE)

clusters <- cluster[["cluster"]]
cluster$cluster <- clusters
movies$cluster <- clusters
```

now we can see which movies belong to each cluster and take a deeper look.

```
#combine clusters with plots_gamma_wider for topic probabilities
plots_gamma_wider <- plots_gamma_wider %>%
  left_join(movies %>% select(Movie.Name, cluster), by = c("document" =
"Movie.Name"))

#take all movies and split them by clusters

cluster_1 <- plots_gamma_wider %>%
  filter(cluster == 1)

cluster_2 <- plots_gamma_wider %>%
  filter(cluster == 2)

cluster_3 <- plots_gamma_wider %>%
  filter(cluster == 3)

cluster_4 <- plots_gamma_wider %>%
  filter(cluster == 4)

cluster_5 <- plots_gamma_wider %>%
  filter(cluster == 5)

cluster_6 <- plots_gamma_wider %>%
  filter(cluster == 6)

cluster_7 <- plots_gamma_wider %>%
  filter(cluster == 7)

cluster_8 <- plots_gamma_wider %>%
  filter(cluster == 8)

cluster_9 <- plots_gamma_wider %>%
  filter(cluster == 9)

cluster_10 <- plots_gamma_wider %>%
  filter(cluster == 10)
```

We can now take the averages of these dataframes to see which topic is associated most with each cluster

#create a function that takes averages of the columns

```
average_columns <- function(df) {  
  # Select only columns named 1 to 20  
  selected_columns <- df %>%  
    select(`1`:`20`)  
  
  # Calculate the column averages  
  column_averages <- colMeans(selected_columns, na.rm = TRUE)  
  
  return(column_averages)  
}
```

use this function for a cluster (1)

```
averages_cluster_1 <- average_columns(cluster_1)  
print(averages_cluster_1)
```

```
##           1           2           3           4           5           6  
## 0.074588920 0.076211741 0.054285215 0.072251828 0.086216236 0.001397285  
##           7           8           9          10          11          12  
## 0.039313256 0.052618655 0.058959002 0.039351455 0.119292812 0.037636649  
##          13          14          15          16          17          18  
## 0.039321689 0.006324736 0.019852370 0.056355101 0.119063649 0.006169812  
##          19          20  
## 0.020930394 0.019859193
```

We can see that these probabilities are pretty small, however a few of them stick out, particularly topics 4 and 14. This indicates that cluster 1 is most associated with topics 4 and 14.

We can now create a word cloud from these topics

```
library(wordcloud)
```

```
## Loading required package: RColorBrewer
```

Define a function to create word clouds for a given topic

```
create_wordcloud <- function(topic_number) {  
  # Filter betas for the chosen topic and select the top words by beta  
  top_words <- betas %>%  
    filter(topic == topic_number) %>%  
    top_n(30, beta) %>% # Adjust 30 to however many top words you want  
    arrange(desc(beta))
```

Generate the word cloud

```
wordcloud(words = top_words$term,  
          freq = top_words$beta,
```




based on these word clouds the genre is not that clear. However we can kind of say it is between western and sci fi

Lets do this again for cluster 6

```
averages_cluster_6 <- average_columns(cluster_6)
print(averages_cluster_6)
```

##	1	2	3	4	5	6
7						
##	0.01356578	0.01664848	0.01663862	0.11156594	0.03579947	0.05687211
	0.01423304					
##	8	9	10	11	12	13
14						
##	0.03471028	0.06257654	0.03771703	0.08897297	0.06540915	0.03172539
	0.10792952					
##	15	16	17	18	19	20
##	0.05843596	0.05231028	0.04244371	0.06560673	0.07168099	0.01515801

this cluster seems to be mostly associated with topics 1 and 12 so let's see their word clouds.

```
create_wordcloud(1)
```

band exercises
father captain
french doctor history
gold takes fortune
gang olcott time mission
truth frail dr war
night party life woman
secret world horse
mysterious
straight danger
american discovers

```
create_wordcloud(12)
```





Joining
these two word clouds it seems that cluster 6 can be classified as crime thriller/mystery movies

Additionally we can look at the highest proportional topics across clusters

```
cluster_topic_averages <- plots_gamma_wider %>%
  group_by(cluster) %>%
  summarize(across(`1`:`20`, mean, na.rm = TRUE)) %>%
  pivot_longer(cols = `1`:`20`, names_to = "topic", values_to =
    "average_gamma")

## Warning: There was 1 warning in `summarize()`.
## [i] In argument: `across(`1`:`20`, mean, na.rm = TRUE)`.
## [i] In group 1: `cluster = 1`.
## Caused by warning:
## ! The `...` argument of `across()` is deprecated as of dplyr 1.1.0.
## Supply arguments directly to `.fns` through an anonymous function instead.
##
## # Previously
##   across(a:b, mean, na.rm = TRUE)
##
## # Now
##   across(a:b, \(x) mean(x, na.rm = TRUE))
```

```
# Plot topic proportions for each cluster
ggplot(cluster_topic_averages, aes(x = factor(topic), y = average_gamma, fill
= factor(cluster))) +
  geom_bar(stat = "identity", position = "dodge") +
  labs(title = "Average Topic Proportions by Cluster", x = "Topic", y =
"Average Gamma") +
  theme_minimal() +
  theme(legend.position = "bottom") +
  scale_fill_brewer(palette = "Set3", name = "Cluster")
```

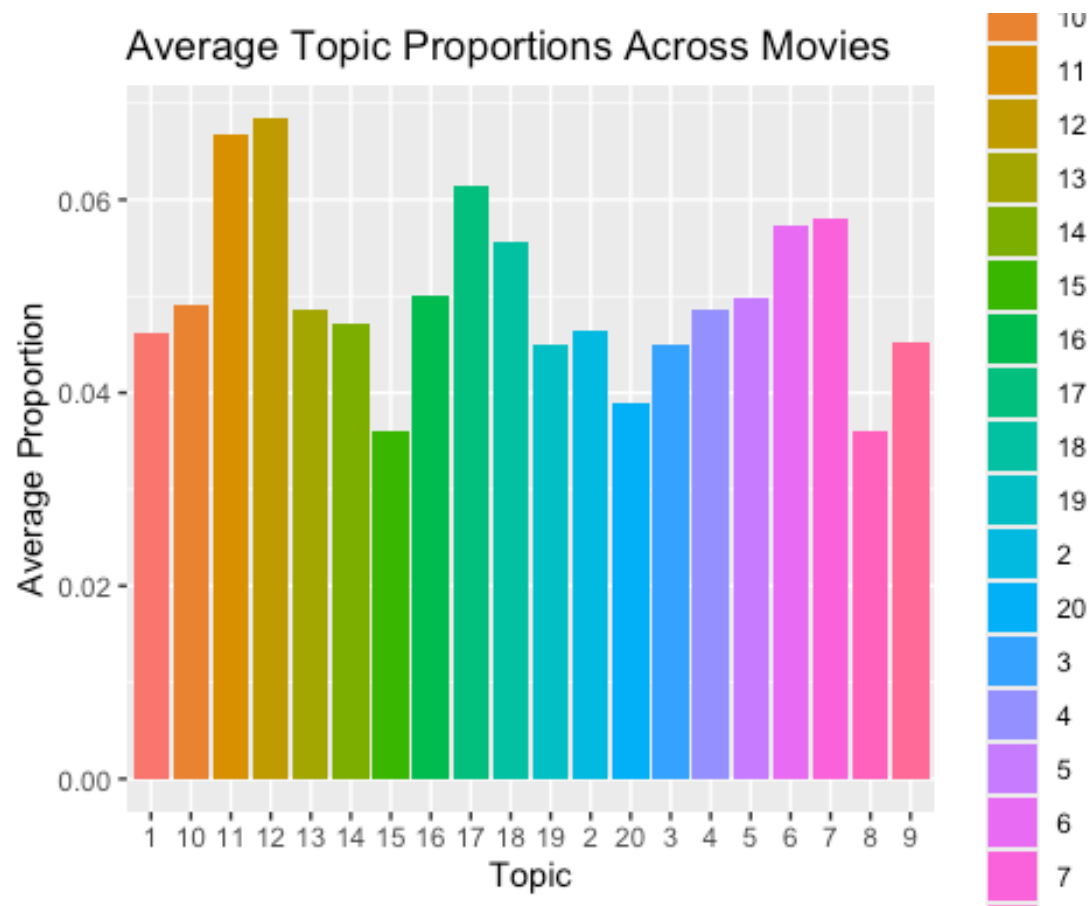


This shows us which clusters pertain the most to each individual topic on a broader level.
(NoteL: I am not sure why topics are unordered)

We can also see the proportion of topics across movies

```
plots_gamma_wider %>%
  pivot_longer(cols = `1`:`20`, names_to = "topic", values_to = "gamma") %>%
  # Pivot all at once
  group_by(topic) %>%
  summarise(avg_proportion = mean(gamma, na.rm = TRUE)) %>%
  ggplot(aes(x = topic, y = avg_proportion, fill = topic)) +
  geom_bar(stat = "identity", position = "dodge") +
```

```
labs(title = "Average Topic Proportions Across Movies", x = "Topic", y = "Average Proportion")
```



This shows the average proportion of each topic across all movies. Which helps to identify which topics are more or less prominent in the movie dataset